

CUDA Acceleration for Edge Detection

ECE 5720

Jingkai Zhang (jz544)
Chen-Tung Chu (cc2396)

Motivation

- Edge detection is commonly implemented in various fields
 - Computer vision
 - Machine learning
- Applications
 - Depth perception
 - Object detection
 - Fingerprint recognition
 - License plate detection
- Computation speed is crucial for such implementation
 - Lower the high-resolution digital images for faster convolution operations.

Edge detection

- What is edge detection?
 - Simplest definition: Sharp changes in the image brightness
 - The points where the image brightness varies sharply are called **the edges** of the image



Source: [Wiki](#)

- Three major edge detection methods
 - Sobel, Canny, Fuzzy logic

Edge detection (cont.)

- How to process the high-resolution digital image?
 - Convolution comes to the rescue

7	2	3	3	8
4	5	3	8	4
3	3	2	8	4
2	8	7	2	7
5	4	4	5	4

*

1	0	-1
1	0	-1
1	0	-1

=

6		

$7 \times 1 + 4 \times 1 + 3 \times 1 + 2 \times 0 + 5 \times 0 + 3 \times 0 + 3 \times -1 + 3 \times -1 + 2 \times -1 = 6$

Source: [medium](#)

- For the purple square:

- *output* =

$$a_{11} * 1 + a_{12} * 0 + a_{13} * (-1) + a_{21} * 1 + a_{22} * 0 + a_{23} * (-1) + a_{31} * 1 + a_{32} * 0 + a_{33} * (-1)$$

a11	a12	a13	a14	a15	a16
a21	a22	a23	a24	a25	a26
a31	a32	a33	a34	a35	a36
a41	a42	a43	a44	a45	a46
a51	a52	a53	a54	a55	a56
a61	a62	a63	a64	a65	a66

6*6 image

*

1	0	-1
1	0	-1
1	0	-1

3*3 filter

=

4*4 image output after edge detection

Source: [GreatLearning](#)

Sobel Edge detection

- Detect the horizontal and vertical edges in images

- Filters:

- $G_x = \begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix} * A$ (vertical edge detection)


- $G_y = \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & 2 & -1 \end{bmatrix} * A$ (horizontal edge detection)

Sobel Edge detection (cont.)

- The Sobel kernels can compute the gradient with smoothing
- $G_x = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} * ([+1 \ 0 \ -1] * A) \quad \Rightarrow \text{increasing the right direction}$
- $G_y = \begin{bmatrix} +1 \\ 0 \\ -1 \end{bmatrix} * ([1 \ 2 \ 1] * A) \quad \Rightarrow \text{increasing the down direction}$
- For each point in the image, the gradient approximation can be combined as the gradient magnitude:
 - $G = \sqrt{G_x^2 + G_y^2}$

Implementation

```
1 function sobel_edge_detector(A)    % A as the image
2   Gx = [-1 0 1; -2 0 2; -1 0 1]
3   Gy = [-1 -2 -1; 0 0 0; 1 2 1]
4
5   image_row = size(A, 1)
6   image_col = size(A, 2)
7   mag = zeros(A)    % Magnitude
8
9   for i=1:image_row-2
10      for j=1:image_col-2
11         S1 = sum(sum(Gx.*A(i:i+2, j:j+2)))
12         S2 = sum(sum(Gy.*A(i:i+2, j:j+2)))
13         mag(i+1, j+1) = sqrt(S1^2 + S2^2)
14      end for
15   end for
16
17   threshold = 70
18   output_image = max(mag, threshold)
19   output_image(output_image == round(threshold)) = 0;
20   return output_image
21 end function
```



Parallel approaches

- OpenMP

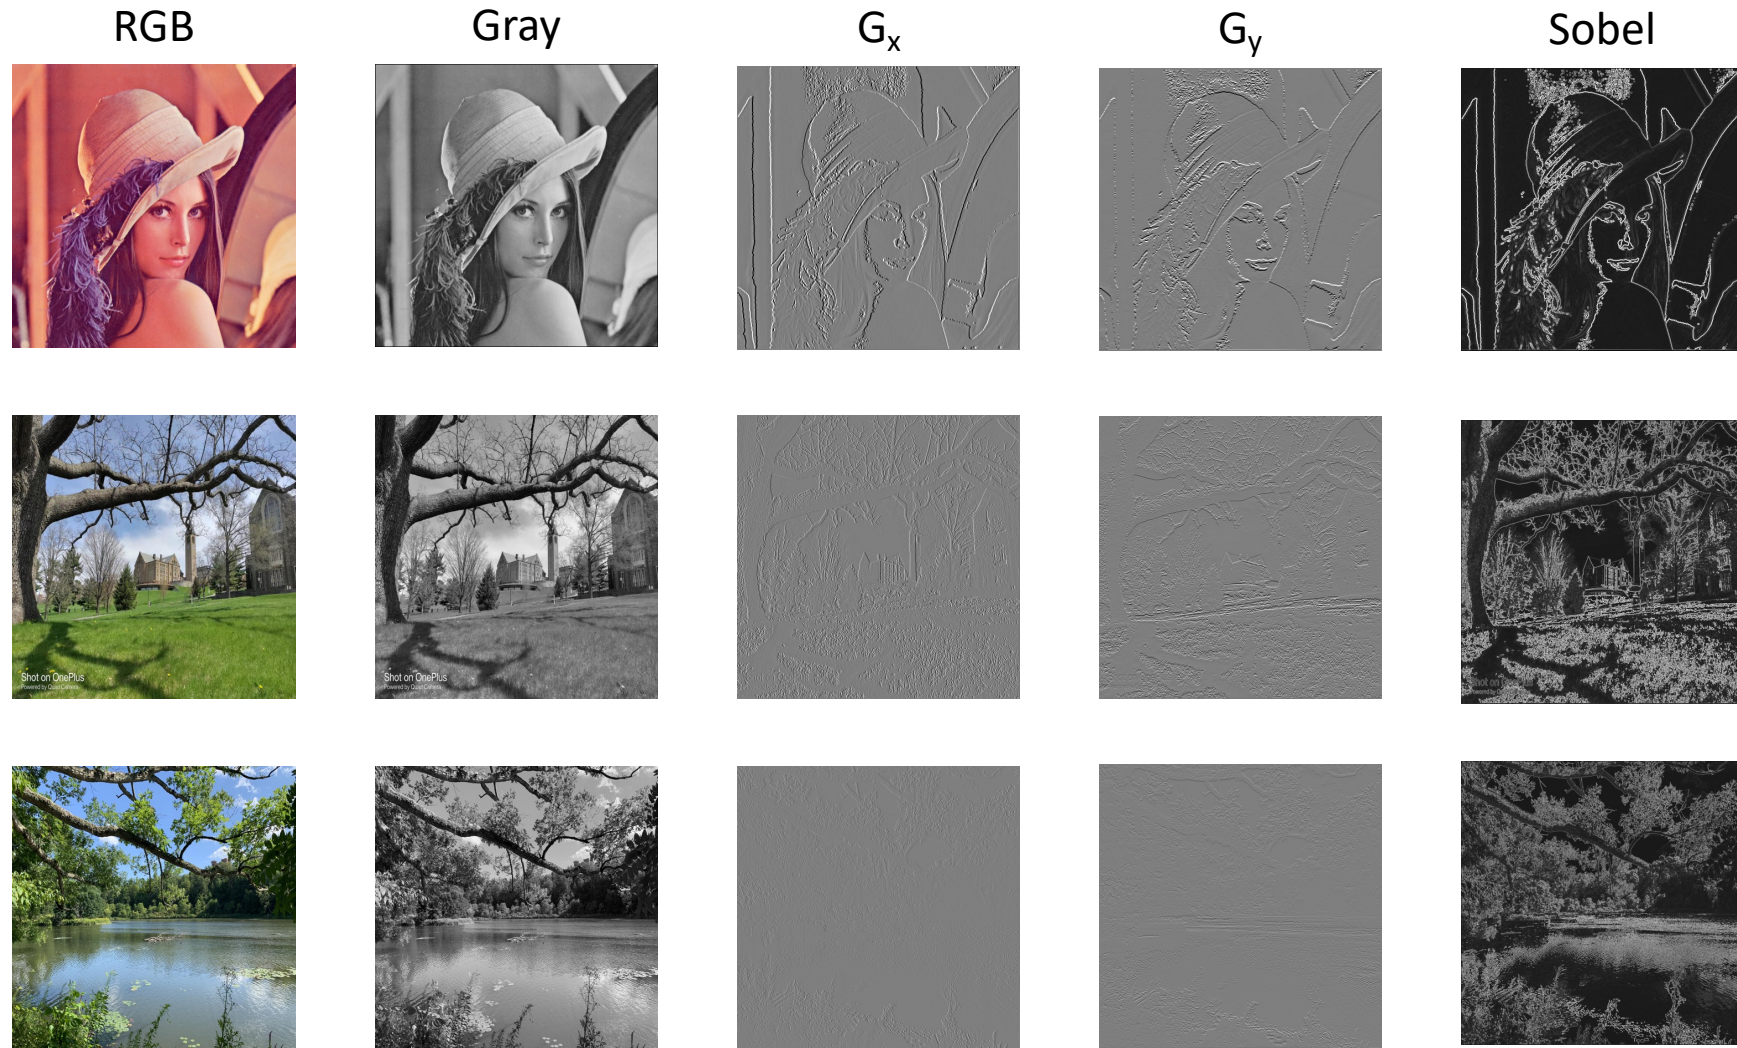
```
1 void sobel_edge_filter_omp(image, out_image){
2   int gx, gy, i, j;
3   #pragma omp parallel for default(shared) private( i, j, gx, gy)
4     for( i = 1; i < image->height - 2; i++ ) {
5       for( j = 1; j < image->width - 2; j++ ) {
6         gx = convolution(i,j,sobel_filter_x);
7         gy = convolution(i,j,sobel_filter_y);
8         out_image->imageData[i][j] = sqrt(gx*gx + gy*gy);
9         out_image->gx[i][j] = gx;
10        out_image->gy[i][j] = gy;
11      }
12    }
13 }
```


Parallel approaches


- Cuda

```
1  __global__ void sobel_edge_filter_cuda(int* dataIn, int* dataOut, int imgHeight, int imgWidth) {
2      int x_index = threadIdx.x + blockIdx.x * blockDim.x;
3      int y_index = threadIdx.y + blockIdx.y * blockDim.y;
4      int index = y_index * imgWidth + x_index;
5      int Gx = 0;
6      int Gy = 0;
7      if (x_index > 0 && x_index < imgWidth - 2 && y_index > 0 && y_index < imgHeight - 2) {
8          Gx = convolution(x_index, y_index, sobel_filter_x);
9          Gy = convolution(x_index, y_index, sobel_filter_y);
10         dataOut[index] = sqrt(Gx * Gx + Gy * Gy);
11     }
12 }
```

Results



Comparisons




Comparisons

images: 100

threads in openMP: 10


Cuda size: grid(30, 30), block(300, 300)



Comparisons (plus loading image time)

SEQUENTIAL

■ Sequential



PARALLEL

■ openMP ■ Cuda

